

NFS Delegations

Table of Contents

Delegations	1
When does an nfs server grant a delegation	1
What does this demonstration show?	2
Example delegation using packet capture	2
Conclusions from the demonstration	8

This is a demonstration of how NFS v4 read delegations operate

Delegations

When the client opens a file, the nfs v4 server has the option of providing a delegation to the client for that particular open file. When the server grants a delegation, the server guarantees the client has full control over the file (according to the parameters of the delegation), and the client assumes management of the file. The client no longer needs to communicate with the server before taking various actions. This greatly reduces the interactions between the client & server, reducing network traffic, and increasing performance.

The server may choose not to grant a delegation, or may recall the delegation at any time, at which point the client must stop assuming full control over the file, and must resume interacting with the server as previously.

When does an nfs server grant a delegation

Speaking for the Linux nfs v4 server in RHEL 7 and newer, the nfs server may grant a read delegation when the client opens a file, the file has not been opened for writing by any clients, and the server is configured to permit file leases.

File leases are configured through the sysctl **fs.leases-enable**. When set to '1' at the time that the nfs client first establishes a session with the server, the client and server will negotiate a backchannel for the server to initiate communications back to the client, if necessary.

```

16 2019-10-12 17:32:18.200013125 192.168.122.60 → 192.168.122.71 NFS 434 V4 Call
(Reply In 18) EXCHANGE_ID
18 2019-10-12 17:32:18.200264649 192.168.122.71 → 192.168.122.60 NFS 286 V4 Reply
(Call In 16) EXCHANGE_ID
21 2019-10-12 17:32:18.200831332 192.168.122.60 → 192.168.122.71 NFS 350 V4 Call
(Reply In 22) CREATE_SESSION
    nfs.create_session.flags.conn_back_chan == 1
22 2019-10-12 17:32:18.201186973 192.168.122.71 → 192.168.122.60 NFS 262 V4 Reply
(Call In 21) CREATE_SESSION
    nfs.create_session.flags.conn_back_chan == 1

```

What does this demonstration show?

This particular demonstration mounts the filesystem using the 'noac' option, which prevents the client from caching file attributes. Ordinarily, the client will cache file attributes, and will perform a check every few seconds to see whether the nfs server's version of each file's attributes has had any updates. Because of this, the client may not know about changes to the file which occur during that small interval. Specifying 'noac' requires the client to perform these checks before every operation, so that attribute changes are detected immediately. This causes a significant performance penalty.

This demonstration shows how that performance penalty can be mitigated when read delegations can be used.

Example delegation using packet capture

The demonstration environment includes 2 nfs clients, both mounting the same nfs export using nfs version 4.1.

```

[root@client1 ~]# mount | grep server
server:/exports on /mnt/server type nfs4
(rw,relatime,sync,vers=4.1,rsize=524288,wsz=524288,namlen=255,acregmin=0,acregmax=0,
acdirmin=0,acdirmax=0,hard,noac,proto=tcp,timeo=600,retrans=2,sec=sys,clientaddr=192.1
68.122.60,local_lock=none,addr=192.168.122.71)

```

```

[root@client2 ~]# mount | grep server
server:/exports on /mnt/server type nfs4
(rw,relatime,sync,vers=4.1,rsize=524288,wsz=524288,namlen=255,acregmin=0,acregmax=0,
acdirmin=0,acdirmax=0,hard,noac,proto=tcp,timeo=600,retrans=2,sec=sys,clientaddr=192.1
68.122.73,local_lock=none,addr=192.168.122.71)

```

on 'client1', open a file read-only; the server grants a read delegation:

```

[root@client1 ~]# exec 10</mnt/server/trace.dat

```

```

121 27.950367616      client1 → server      NFS 350 V4 Call (Reply In 122) SEQUENCE
| PUTFH | OPEN DH: 0x8e2272a5/trace.dat | GETFH | ACCESS FH: 0x00000000, [Check: RD MD
XT XE] | GETATTR
      rpc.xid == 0x75438f20
      nfs.fh.hash == 0xde61574d
      nfs.fh.hash == 0x8e2272a5
nfs.open4.share_access == OPEN4_SHARE_ACCESS_READ
nfs.open4.share_deny == OPEN4_SHARE_DENY_NONE
nfs.open.claim_type == CLAIM_NULL
nfs.open.opentype == OPEN4_NOCREATE
122 27.950603133      server → client1      NFS 474 V4 Reply (Call In 121) SEQUENCE
| PUTFH | OPEN StateID: 0xc262 | GETFH | ACCESS, [Access Denied: XE], [Allowed: RD MD
XT] | GETATTR
      rpc.xid == 0x75438f20
      nfs.fh.hash == 0x8e2272a5
      nfs.fh.hash == 0xde61574d
nfs.open.delegation_type == OPEN_DELEGATE_READ

```

on 'client1', run 'dd' to open a file, read its contents, and close the file again:

```

[root@client1 ~]# dd if=/mnt/server/trace.dat of=/dev/null
14128+1 records in
14128+1 records out
7233628 bytes (7.2 MB) copied, 0.0870879 s, 83.1 MB/s

```

Because the nfs client already had a read delegation, and because 'dd' was opening the file read-only again, the client did not need to communicate with the server in order to do the file 'OPEN' or 'CLOSE' operations when opening or closing the file. Only the 'READ' calls and replies were exchanged between the client and server. In addition, due to the delegation, the client did not need to check the file attributes prior to each read.

on 'client2', open the file read-only:

```

[root@client2 ~]# exec 10</mnt/server/trace.dat

```

the nfs server also grants a read delegation to 'client2':

```

1387 63.374995537      client2 → server      NFS 350 V4 Call (Reply In 1388)
SEQUENCE | PUTFH | OPEN DH: 0x8e2272a5/trace.dat | GETFH | ACCESS FH: 0x00000000,
[Check: RD MD XT XE] | GETATTR
      rpc.xid == 0xfb8f09da
      nfs.fh.hash == 0xde61574d
      nfs.fh.hash == 0x8e2272a5
nfs.open4.share_access == OPEN4_SHARE_ACCESS_READ
nfs.open4.share_deny == OPEN4_SHARE_DENY_NONE
nfs.open.claim_type == CLAIM_NULL
nfs.open.opentype == OPEN4_NOCREATE
1388 63.375491192      server → client2      NFS 474 V4 Reply (Call In 1387)
SEQUENCE | PUTFH | OPEN StateID: 0xdfd3 | GETFH | ACCESS, [Access Denied: XE],
[Allowed: RD MD XT] | GETATTR
      rpc.xid == 0xfb8f09da
      nfs.fh.hash == 0x8e2272a5
      nfs.fh.hash == 0xde61574d
nfs.open.delegation_type == OPEN_DELEGATE_READ

```

run 'dd' on 'client2' as well:

```

[root@client2 ~]# dd if=/mnt/server/trace.dat of=/dev/null
14128+1 records in
14128+1 records out
7233628 bytes (7.2 MB) copied, 0.12299 s, 58.8 MB/s

```

again, because the client had a read delegation, the 'OPEN' and 'CLOSE' operations were not communicated with the server, and the client did not check file attributes prior to each 'READ' operation.

Now, on 'client2', open the file in a read-write mode:

```

[root@client2 ~]# exec 11<>/mnt/server/trace.dat

```

Because the client knows that opening the file read-write will conflict with its read delegation, the client calls 'DELEGRETURN' to inform the server that it will no longer use the delegation:

```

1819 81.896349549      client2 → server      NFS 290 V4 Call (Reply In 1820)
SEQUENCE | PUTFH | GETATTR FH: 0xde61574d | DELEGRETURN StateID: 0xc368
      rpc.xid == 0x0d9009da
      nfs.fh.hash == 0xde61574d
1820 81.896668788      server → client2      NFS 230 V4 Reply (Call In 1819)
SEQUENCE | PUTFH | GETATTR | DELEGRETURN
      rpc.xid == 0x0d9009da
      nfs.fh.hash == 0xde61574d

```

After returning the delegation, 'client2' then makes an 'OPEN' call to open the file in read-write

mode:

```
1821 81.897544067      client2 → server      NFS 338 V4 Call (Reply In 1822)
SEQUENCE | PUTFH | OPEN DH: 0xde61574d/ | ACCESS FH: 0xde61574d, [Check: RD MD XT XE]
| GETATTR FH: 0xde61574d
      rpc.xid == 0x0e9009da
      nfs.fh.hash == 0xde61574d
nfs.open4.share_access == OPEN4_SHARE_ACCESS_BOTH
nfs.open4.share_deny == OPEN4_SHARE_DENY_NONE
nfs.open.claim_type == CLAIM_FH
nfs.open.opentype == OPEN4_NOCREATE
```

Because the nfs server has also given a read delegation to 'client1', the server must first recall the delegation from 'client1' before allowing 'client2' to open the file with a mode that permits writing. The server therefore responds to the 'OPEN' call with 'NFS4ERR_DELAY' to instruct 'client2' to wait briefly before making the call again:

```
1822 81.897906741      server → client2      NFS 166 V4 Reply (Call In 1821)
SEQUENCE | PUTFH | OPEN Status: NFS4ERR_DELAY
      rpc.xid == 0x0e9009da
      nfs.fh.hash == 0xde61574d
```

The server next informs 'client1' that it must return its read delegation. For this, it makes a 'CB_RECALL' call using the rpc program named 'NFS_CB', rather than making a call using the 'NFS' program. The nfs client acknowledges that it has received the message that it must return the delegation:

```
1823 81.898057606      server → client1      NFS CB 278 V1 CB_COMPOUND Call (Reply
In 1824) <EMPTY> CB_SEQUENCE;CB_RECALL
      rpc.xid == 0xdd4fbbb4
      nfs.fh.hash == 0xde61574d
1824 81.898703209      client1 → server      NFS CB 154 V1 CB_COMPOUND Reply (Call
In 1823) <EMPTY> CB_SEQUENCE;CB_RECALL
      rpc.xid == 0xdd4fbbb4
      nfs.fh.hash == 0xde61574d
```

Note that the nfs callback is initiating a call back to the client, not responding to a call from the client; when using nfs v4.0, this callback will over a separate tcp port, negotiated at the time that the session was initially set up. Beginning with nfs v4.1, the existing tcp port 2049 connection is reused for the backchannel communication.

As instructed, 'client1' now returns its read delegation:

```

1825 81.898707139      client1 → server      NFS 290 V4 Call (Reply In 1827)
SEQUENCE | PUTFH | GETATTR FH: 0xde61574d | DELEGRETURN StateID: 0xded9
      rpc.xid == 0x89438f20
      nfs.fh.hash == 0xde61574d
1827 81.899046375      server → client1      NFS 230 V4 Reply (Call In 1825)
SEQUENCE | PUTFH | GETATTR | DELEGRETURN
      rpc.xid == 0x89438f20
      nfs.fh.hash == 0xde61574d

```

After 'client2' received the 'NFS4ERR_DELAY' response, it slept briefly. It now reattempts the 'OPEN' call to open the file read-write:

```

1830 82.001078714      client2 → server      NFS 338 V4 Call (Reply In 1831)
SEQUENCE | PUTFH | OPEN DH: 0xde61574d/ | ACCESS FH: 0xde61574d, [Check: RD MD XT XE]
| GETATTR FH: 0xde61574d
      rpc.xid == 0x0f9009da
      nfs.fh.hash == 0xde61574d
      nfs.open4.share_access == OPEN4_SHARE_ACCESS_BOTH
      nfs.open4.share_deny == OPEN4_SHARE_DENY_NONE
      nfs.open.claim_type == CLAIM_FH
      nfs.open.opentype == OPEN4_NOCREATE
1831 82.001388279      server → client2      NFS 386 V4 Reply (Call In 1830)
SEQUENCE | PUTFH | OPEN StateID: 0xd124 | ACCESS, [Access Denied: XE], [Allowed: RD MD
XT] | GETATTR
      rpc.xid == 0x0f9009da
      nfs.fh.hash == 0xde61574d
      nfs.open.delegation_type == OPEN_DELEGATE_NONE

```

The server responded that the 'OPEN' was successful, but gives no delegation to the client.

Because it no longer has a delegation in place, when 'client1' runs the 'dd' command again, the 'dd' progresses much more slowly, since the client has to regularly verify that the file attributes have not changed:

```

[root@client1 ~]# dd if=/mnt/server/trace.dat of=/dev/null
14128+1 records in
14128+1 records out
7233628 bytes (7.2 MB) copied, 4.66734 s, 1.5 MB/s

```

When 'client2' closes the read-write file, the client simply reduces its access level to read-only, using an 'OPEN_DOWNGRADE' call:

```

30105 120.627189249 client2 → server NFS 286 V4 Call (Reply In 30106)
SEQUENCE | PUTFH | OPEN_DOWNGRADE
      rpc.xid == 0x109009da
      nfs.fh.hash == 0xde61574d
      nfs.open4.share_access == OPEN4_SHARE_ACCESS_READ
      nfs.open4.share_deny == OPEN4_SHARE_DENY_NONE
30106 120.627525362 server → client2 NFS 182 V4 Reply (Call In 30105)
SEQUENCE | PUTFH | OPEN_DOWNGRADE
      rpc.xid == 0x109009da
      nfs.fh.hash == 0xde61574d

```

Note that neither 'client1' nor 'client2' is given a read delegation again; this is done at the time of an 'OPEN' call.

on 'client2', close the file which was open read-only:

```
[root@client2 ~]# exec 10>&-
```

```

30108 124.662539565 client2 → server NFS 294 V4 Call (Reply In 30109)
SEQUENCE | PUTFH | GETATTR FH: 0xde61574d | CLOSE StateID: 0xd489
      rpc.xid == 0x119009da
      nfs.fh.hash == 0xde61574d

30109 124.662875472 server → client2 NFS 246 V4 Reply (Call In 30108)
SEQUENCE | PUTFH | GETATTR | CLOSE
      rpc.xid == 0x119009da
      nfs.fh.hash == 0xde61574d

```

on 'client1', close the file as well:

```
[root@client1 ~]# exec 10>&-
```

```

58386 146.150172048 client1 → server NFS 294 V4 Call (Reply In 58387)
SEQUENCE | PUTFH | CLOSE StateID: 0xc262 | GETATTR FH: 0xde61574d
      rpc.xid == 0xf6b18f20
      nfs.fh.hash == 0xde61574d
58387 146.150582719 server → client1 NFS 246 V4 Reply (Call In 58386)
SEQUENCE | PUTFH | CLOSE | GETATTR
      rpc.xid == 0xf6b18f20
      nfs.fh.hash == 0xde61574d

```

On 'client1', run the 'dd' again:

```
[root@client1 ~ ]# dd if=/mnt/server/trace.dat of=/dev/null
14128+1 records in
14128+1 records out
7233628 bytes (7.2 MB) copied, 0.0465882 s, 155 MB/s
```

```
58398 151.263072172 client1 → server NFS 338 V4 Call (Reply In 58399)
SEQUENCE | PUTFH | OPEN DH: 0xde61574d/ | ACCESS FH: 0xde61574d, [Check: RD MD XT XE]
| GETATTR FH: 0xde61574d
    rpc.xid == 0xfbb18f20
    nfs.fh.hash == 0xde61574d
    nfs.open4.share_access == OPEN4_SHARE_ACCESS_READ
    nfs.open4.share_deny == OPEN4_SHARE_DENY_NONE
    nfs.open.claim_type == CLAIM_FH
    nfs.open.opentype == OPEN4_NOCREATE
58399 151.263419750 server → client1 NFS 422 V4 Reply (Call In 58398)
SEQUENCE | PUTFH | OPEN StateID: 0xfbb14 | ACCESS, [Access Denied: XE], [Allowed: RD MD
XT] | GETATTR
    rpc.xid == 0xfbb18f20
    nfs.fh.hash == 0xde61574d
    nfs.open.delegation_type == OPEN_DELEGATE_READ
```

Note that the server has granted a read delegation again, and the file read progresses quickly again.

Because there are no additional file handles open for the file, 'client1' closes the file immediately:

```
58401 151.309305617 client1 → server NFS 278 V4 Call (Reply In 58402)
SEQUENCE | PUTFH | CLOSE StateID: 0xfbb14
    rpc.xid == 0xfcb18f20
    nfs.fh.hash == 0xde61574d
58402 151.309690377 server → client1 NFS 182 V4 Reply (Call In 58401)
SEQUENCE | PUTFH | CLOSE
    rpc.xid == 0xfcb18f20
    nfs.fh.hash == 0xde61574d
```

Conclusions from the demonstration

When an nfs client is granted a read delegation, it has a guarantee from the server that the file will not change; it is therefore unnecessary for the client to perform any checking of the file attributes. This dramatically improves the client performance.

After the client had to return its delegation, it needed to constantly check the file attributes, significantly reducing the overall transfer rate.

Note that the benefits provided by a read delegation are not always this dramatic. Mounting with 'noac' shows an extreme, worst-case scenario.